


02: Mesh Data Structures

Dr. Hamish Carr


Data Structures

- Based on identifying the requirements:
 - topological – mesh assumptions
 - algorithmic – what operations we will need
 - resource – how much memory / bandwidth
 - complexity – preprocessing / runtime cost
- Throughout, we will assume IEEE floats (4B)



Topological Requirements

- Meshes are manifold (if possible)
- If not, the boundary is clearly marked
- Meshes are triangulations
 - Easily enforced – subdivide polygons
- Sometimes generalise to quads
- Almost never use larger mesh elements




Algorithmic Requirements

- Operations will include:
 - Rendering
 - GPU transfer
 - Iteration for mesh processing
 - Insertion / deletion / modification

Operations

- a. access to vertices, edges, faces
- b. iteration over vertices, edges, faces (any order)
- c. access to endpoints of edge
- d. access to faces for each edge
- e. oriented traversal of edges (walk around face)
- f. oriented traversal of vertices (ditto) c,e
- g. oriented traversal of edges (around vertex) d,e
- h. oriented traversal of faces (ditto) d,g




Example: Smooth Normals

- We'll use this as a running example
- For each vertex
 - *average* the adjacent normals
 - to produce an approximate vertex normal
- Strictly speaking, don't need to iterate around
 - We just need to iterate over adjacent faces
 - Any order will do

Resource Requirements

- Memory footprint is the key issue
- It also influences bandwidth
 - Memory transfer to GPU for rendering
 - Not all operations are on GPU
- Ideal language these days is C++
 - Because it enforces the POD rule




Runtime Polymorphism

- OO languages have *runtime* polymorphism
 - decides which function to call at runtime
 - needs to store information on object class
 - typically done with trap table pointer
 - pointer to details of class in memory
 - object name, method base addresses, &c.


Trap Table Penalty

- This trap table pointer is on CPU
 - it's *totally* useless on GPU
- If you have a class representing a point
 - `Point3 { float x, y, z; NULL *trapTablePtr; }`
 - 12B for your payload, 4B or 8B for the pointer
- GPU bandwidth goes up by 33% or 67%
- Unless you strip it out first (i.e. copy!)



POD Loophole

- If you have no runtime polymorphism
 - i.e. no virtual methods *at all*
 - you don't need the trap table pointer
- C++ guarantees not to create one
 - So you can copy your data straight to GPU
- And arrays are even better!
 - One memory copy for the entire array



Complexity Requirements


- Algorithms should be $O(n)$ if possible
- Insertions / deletions should be $O(1)$ cost
- Iterations should also be $O(1)$ cost
- $O(n \lg n)$ should be reserved for preprocessing
- $O(\lg n)$ is the maximum runtime lookup cost

Mesh Data Structures

- Face
- Indexed Face
- Winged-Edge
- Half-Edge
- Directed Edge

Face Data Structure

- Simplest approach:
 - Each face consists of 3 (xyz) points
 - Render uses `glVertex()` or equivalent
 - Most operations iterate over faces
 - Cannot (usually) enforce constraints
 - Attributes (normals, &c.) need to be repeated
- Colloquially called *triangle/polygon soup*



Face Operations

- Reading / rendering is easy (but not fast)
- Vertices stored 6x, prone to matching errors
- Cannot iterate around vertices
- Cannot (easily) cross edge to next face
- Can iterate over vertices 6x, edges 2x
 - by iterating through faces
- Can iterate over, around faces
- No per vertex/edge storage

Smooth Normals- Face







```
// FACE DATA STRUCTURE
for (each face f)
{ // loop through faces
  // start with a zero vector
  normal = zero
  for (each vertex v on face f)
  { // loop through vertices on face
    // loop through all faces (including this one)
    for (each face f1)
    { // loop through faces again
      for (each vertex v1 on face f1)
      { // loop through vertices on faces
        // if it's the same vertex (look out for roundoff errors!)
        if (v1 == v)
        { // vertex match
          compute normal n_f1 for face f1
          normal += n_f1
        } // vertex match
      } // loop through vertices on faces
    } // loop through faces again
  }
  // note: we don't need to keep the degree or divide through
  // because we're going to normalise anyway!
  normalise(normal)
} // loop through faces
```

- Hopelessly inefficient – $O(v^2)$
- Recomputes each normal six times




(Old) Face Improvements

- Display lists – save bandwidth, cache on GPU
- Implicit vertex reuse
- Line Strips & Loops
- Triangle Fans & Strips
- Quads, Quad Strips
- Polygons




Lines, Strips & Loops




GL_LINES:

v0v1
v1v2
v2v3
v3v4
v4v5
v5v0
2 v / e



GL_LINE_STRIP:


v0v1
(v1)v2
(v2)v3
(v3)v4
(v4)v5
~ 1 v / e



GL_LINE_LOOP:

v0v1
(v1)v2
(v2)v3
(v3)v4
(v4)v5
(v5)(v0)
~ 1 v / e

Triangle Fans/Strips



GL_TRIANGLE_FAN:

$v_0v_1v_2$


$(v_0)(v_2)v_3$

$(v_0)(v_3)v_4$

$(v_0)(v_4)v_5$

$\sim 1.16 \text{ v/e}$

(degree 6 vertex)



GL_TRIANGLE_STRIP:

$v_0v_1v_2$

$(v_2)(v_1)v_3$

$(v_2)(v_3)v_4$

$(v_4)(v_3)v_5$

$\sim 1 \text{ v/e}$

(for long strips)


alternate CW/CCW

NP-hard to find them

easier on regular grids


Face Costs

- Memory / bandwidth cost:
 - $3 \times 3 \times 4\text{B}$ per face = 36B
 - $f \sim 2v$, so 72B per vertex
- Every `glVertex()` call has overhead
 - so no advantage for batched operations
- Horrendously inefficient operations
- But easy and lowest common denominator



Indexed Face Data Structure

- Store a *single* array of vertices / attributes
- Store their indices (or pointers) for each face
 - Pointers are no longer cheaper
 - *Most* PUs have 1-op array lookup
- Render uses vertex arrays / VBOs
- But operations aren't so easy




Vertex Arrays

```
GLint vertices[8][3] = {  -1, -1, -1,   -1, -1,  1,   -1,  1, -1,   -1,  1,  1,
                          1, -1, -1,    1, -1,  1,    1,  1, -1,    1,  1,  1};
GLint normals[6][3] =   {  -1,  0,  0,    0, -1,  0,    0,  0, -1,
                          1,  0,  0,    0,  1,  0,    0,  0,  1};
unsigned int triangles[36] = {  0, 1, 3,   0, 3, 2,   0, 1, 4,   1, 4, 5,
                              0, 2, 4,   2, 6, 4,   5, 4, 6,   5, 6, 7,
                              2, 3, 7,   2, 7, 6,   1, 5, 7,   1, 7, 3};

glEnableClientState(GL_VERTEX_ARRAY);
glVertexPointer(3, GL_INT, 0, vertices);
glEnableClientState(GL_NORMAL_ARRAY);
glNormalPointer(3, GL_INT, 0, normals);

// this runs it's own loop internally
glDrawElements(GL_TRIANGLES, 36, GL_UNSIGNED_INT, indices);

// turn it off when we're done
glDisableClientState(GL_VERTEX_ARRAY);
glDisableClientState(GL_NORMAL_ARRAY);
```



- Replace individual glVertex() calls
- Pass *entire arrays* to library
- Push loop through vertices into library call



Vertex Buffer Objects

- Same idea, but data transferred to VRAM
- I.e. we cache the data & save the bandwidth
- Which means memory management:
 - allocation: `glGenBuffers`, `glDeleteBuffers`
 - binding: `glBindBuffer`
 - transfer: `glBufferData`
- And more is moving onto GPU with Vulkan



Indexed Face Operations

- Reading / rendering is easy & fast
- Cannot iterate around vertices
 - But can iterate over them
- Cannot (easily) swap faces across edge
- Can iterate over, around faces
- No per edge storage
- Overall weakness – edge operations



Smooth Normals – Indexed Face

```
// INDEXED FACE DATA STRUCTURE
// initialise the normals to zeros for each vertex
for (each vertex v)
    normal[v] = zero

// now loop through the faces
for (each face f)
{ // loop through faces
  // start with a zero vector
  compute normal n_f for face f
  // loop through face's vertices
  for (each vertex v on face f)
  { // per vertex
    // add the face normal to the vertex normal
    normal[v] += n_f
  } // per vertex
} // loop through faces

// now normalise them all
for (each vertex v)
    normalise(normal[v])
```

- Much better – $O(v)$ computation
- But it cannot iterate *around* vertices yet



Reading Polygon Soup

```
int main(int argc, char **argv)
{ // main()
  std::vector<Cartesian3> raw_vertices;
  while (std::cin.good())
  { // loop to read
    Cartesian3 next;
    std::cin >> next.x >> next.y >> next.z;
    raw_vertices.push_back(next);
  } // loop to read

  // now we loop to set vertex IDs for each existing vertex
  std::vector<long> vertexID(raw_vertices.size(), -1);

  // set the initial vertex ID
  long nextVertexID = 0;

  // loop through the vertices
  for (long vertex = 0; vertex < raw_vertices.size(); vertex++)
  { // vertex loop
    // first see if the vertex already exists
    for (long other = 0; other < vertex; other++)
    { // per other
      if (raw_vertices[vertex] == raw_vertices[other])
        vertexID[vertex] = vertexID[other];
    } // per other
    // if not set, set to next available
    if (vertexID[vertex] == -1)
      vertexID[vertex] = nextVertexID++;
  } // vertex loop
```



Writing Indexed Faces

```
// the first four lines will be skipped completely
std::cout << "#" << std::endl;
std::cout << "# Created for Leeds COMP 5821M Autumn 2019" << std::endl;
std::cout << "#" << std::endl;
std::cout << "#" << std::endl;
std::cout << "# Surface vertices=" << nextVertexID << " faces="
    << raw_vertices.size()/3 << std::endl;
std::cout << "#" << std::endl;

// id of next vertex to write
long writeID = 0;


// loop to write the vertices out
for (long vertex = 0; vertex < raw_vertices.size(); vertex++)
{ // per vertex
    // if it's the first time found
    if (writeID == vertexID[vertex])
    { // first time found
        // print out the coordinates
        std::cout << "Vertex " << writeID << " " << raw_vertices[vertex].x << " "
            << raw_vertices[vertex].y << " " << raw_vertices[vertex].z << std::endl;
        // increment the ID number
        writeID++;
    } // first time found
} // per vertex

// loop to write the face vertices out
for (long face = 0; face < raw_vertices.size()/3; face++)
{ // per face
    std::cout << "Face " << face << " " << vertexID[3 * face] << " "
        << vertexID[3 * face + 2] << " " << vertexID[3 * face + 1] << std::endl;
} // per face

// done
return 0;
} // main()
```


Indexed Face Costs

- Memory / bandwidth cost:
 - $3 \times 4\text{B}$ per vertex = 12B
 - $f \sim 2v$, so $6 \times 4\text{B}$ indices per vertex
 - Total: 36B / vertex
- Efficient rendering
- Default in practice: OBJ files / VBOs
- But limited in its operations




Improvements Needed

- Mesh *processing* requires more operations
- In particular, iterating in *cyclic* order
- This allows us to preserve topology
 - e.g. testing 2-manifold conditions
 - can't (easily) do this with indexed face
- We will need to track vertices, faces, edges
 - Each will need index + attribute storage



Face-Based Structures

- Each vertex stores:
 - position 12B
 - 1 face index (first) 4B 16B/vert
- Each face stores:
 - 3 vertex indices 12B
 - 3 neighbours (face indices) 12B 24B/face
 - (2 faces / vertex) 48B /vert
- Total: 64B / vert



Smooth Normals – Face-Based

```
// FACE-BASED DATA STRUCTURE
// loop through vertices
for (each vertex v)
{ // for each vertex
  // initialise normal to zero
  normal[v] = zero
  // start loop control variable at first face
  face = first[v]
  do { // do loop around neighbours
    compute normal n_f for face f
    normal[v] += n_f
    // now find which neighbour to use
    if (faceV[0] == v)
      face = neighbour[0]
    else if (faceV[1] == v)
      face = neighbour[1]
    else if (faceV[2] == v)
      face = neighbour[2]
  } // do loop around neighbours
  while (face != first[v])

  // now normalise
  normalise(normal[v])
} // for each vertex
```


- Iteration around vertices now possible
- But lots of if statements
- Starts failing with non-triangular faces



Edge-Based Structures

- These start *explicitly* representing edges
- And using them to store cyclic edge loops
- Each edge points to the next/previous edge
- At each end, and for each face
- Paradigm is the *winged edge* (Baumgart, 1972)

Winged Edge



Winged Edge Structures

- Each vertex stores:
 - position 12B
 - 1 edge index (first) 4B16B/vertex
- Each edge stores:
 - 2 vertex indices 8B
 - 2 face indices 8B
 - 2 edge indices (next) 8B
 - 2 edge indices (prev) 8B32B/edge
- (3 edges / vertex) 96B/vertex
- Each face stores:
 - 1 edge index (first) 4B8B /vert
- (2 faces / vertex)
- Total: 120B / vert




Smooth Normals – Winged Edge

```
// WINGED-EDGE DATA STRUCTURE
// loop through vertices
for (each vertex v)
{ // for each vertex
  // initialise normal to zero
  normal[v] = zero
  // start loop control variable at first edge
  edge = firstEdge[v]
  do { // do loop around neighbours
    // retrieve the corresponding face
    face = edge.face[0]
    // note: each vertex will need retrieval through edge
    // there are some further optimisations available here
    compute normal n_f for face f
    normal[v] += n_f
    // now step forward
    if (edge == firstEdge[end[0]])
      edge = edge.next[0]
    else
      edge = edge.next[1]
  } // do loop around neighbours
  while (edge != firstEdge[v])

  // now normalise
  normalise(normal[v])
} // for each vertex
```

- This is cleaner, and handles arbitrary meshes
- Face can store vertex IDs as well (+24B/v)
- Still need if statements to track things



Half-Edge Structures

- Break winged-edges into a pair of half-edges
 - Mantyla 1988, Kettner 1999
- Each half edge has:
 - vertex, face, next, prev
 - ID of paired half-edge
 - But store them in pairs
 - Use XOR on bit 0 to flip between them
- Same cost as winged-edge, simpler code

Smooth Normals – Half Edge

```
// HALF-EDGE DATA STRUCTURE
// loop through vertices
for (each vertex v)
{ // for each vertex
  // initialise normal to zero
  normal[v] = zero
  // start loop control variable at first edge
  halfedge = firstHalfEdge[v]
  do { // do loop around neighbours
    // retrieve the corresponding face
    face = halfedge.face[0]
    // note: each vertex will need retrieval through edge
    // there are some further optimisations available here
    compute normal n_f for face f
    normal[v] += n_f
    // now step forward
    edge = edge.next
  } // do loop around neighbours
  while (edge != firstHalfEdge[v])

  // now normalise
  normalise(normal[v])
} // for each vertex
```

- This is nice, but still memory-inefficient
- We need to reduce the per edge costs



Directed Edge Structures

- Campagna 1998
- Stores the directed-edges *per triangle*
- I.e. in groups of three
- Assumes arithmetic cheap, memory not
- Uses modular arithmetic ($\%3$) instead of XOR
- And uses it for implicit storage of topology



Conversion

- Each directed edge d is no. i on face f :
 - $d = 3f + i$ $f = h/3$ $i = h \% 3$
 - d is directed face index
 - f is face index
 - i is 0, 1, or 2 in CCW order
- $\text{next} = 3f + (i + 1) \% 3$
- $\text{prev} = 3f + (i + 2) \% 3$



Smooth Normals – Directed Edge

```
// DIRECTED-EDGE DATA STRUCTURE
// loop through vertices
for (each vertex v)
{ // for each vertex
  // initialise normal to zero
  normal[v] = zero
  // start loop control variable at first edge
  halfedge = firstHalfEdge[v]
  do { // do loop around neighbours
    // retrieve the corresponding face
    face = halfedge / 3
    whichEdge = halfedge % 3
    // note: each vertex will need retrieval through edge
    // there are some further optimisations available here
    compute normal n_f for face f
    normal[v] += n_f
    edge = 3 * face + (whichEdge + 1) % 3
  } // do loop around neighbours
  while (edge != firstHalfEdge[v])

  // now normalise
  normalise(normal[v])
} // for each vertex
```

- Right, let's look at the cost




Directed Edge Structures

- Each vertex stores:
 - position 12B
 - 1 edge index (first) 4B 16B/vertex
- Each directed edge stores:
 - 1 vertex index 4B
 - 1 paired directed edge index 4B 8B/directed edge
 - (6 directed edges / vertex) 48B / vertex
- Each face stores:
 - nothing
- Total: 64B/ vertex



Implementation

- We will do this C-style, not OO
- so we get that efficiency as well
- the vertices are stored in a vertex array
- and the faces have vertex indices in an array
 - since directed edges are stored per face
 - their vertex indices are *also* the face's list



Code (Finally)

```
// data structure for directed-edge processing
class Mesh
{ // Mesh
public:
    std::vector<Point3D> position;           // xyz
    std::vector<indexType> firstDirectedEdge; // per vertex
    std::vector<indexType> faceVertices;     // doubles as "directed edge to"
    std::vector<indexType> otherHalf;        // pairs the directed edges

    void Render()
    { // Render()
        // use the arrays for vertex array calls
        // VBO's are similar, but a bit messier
        glEnableClientState(GL_VERTEX_ARRAY);
        glVertexPointer(3, GL_INT, 0, position);
        glDrawElements(GL_TRIANGLES, faceVertices.size(), GL_UNSIGNED_INT, faceVertices);
        glDisableClientState(GL_VERTEX_ARRAY);
    } // Render()
}; // Mesh
```

- After all that mess, it's surprisingly simple
- But it's not perfect



Directed Edge Comments

- A boundary edge does not have a pair
 - So store a negative index
- Similar structure for quads is easy to construct
 - but we can't mix the two
- Libraries exist (OpenMesh, CGAL, &c.)
- But most people roll their own
 - Because their *attributes* vary

